

chapter 5 Performance Analysis

2021/6/27

Regret Analysis

- A player chooses an action $\theta^{(t)} \in K$ every t period, where K is a feasible set of actions.
- The cost function $f^{(t)}$ determines the cost $f^{(t)}(\theta^{(t)})$ for action $\theta^{(t)}$.
- The player decides his action based on the strategy.

Regret Analysis

- How does the player choose an action which minimizes a total cost $\sum f^{(t)}(\theta^{(t)})$?
- Can the cost function be minimized even if it is not unknown?
- We introduce a regret about the strategy.

definition (Regret)

The difference between the total cost of an action based on a strategy A and the total cost of the optimal strategy θ^ is defined as the regret $\text{Regret}(A)$ of strategy A .*

$$\text{Regret}(A) = \sum_{t=1}^T f^{(t)}(\theta^{(t)}) - \sum_{t=1}^T f^{(t)}(\theta^*)$$

Regret Analysis

Regret analysis in online learning

- Let action be the parameter of the online learner $\boldsymbol{\theta}^{(t)} \in \mathbb{R}^m$ given the training data $(\boldsymbol{x}^{(t)}, y^{(t)})$.
- Let the cost function be a loss function $f^{(t)} = (\boldsymbol{x}^{(t)}, y^{(t)}, \boldsymbol{\theta})$.
- In this case, the optimal strategy is the strategy that chooses the action that minimizes the cost function for all training data.

Follow the Leader

- At first, We consider strategy for choosing action that minimizes the total cost to date.

$$\boldsymbol{\theta}^{(t)} = \arg \min_{\boldsymbol{\theta} \in K} \sum_{i=1}^{t-1} f^{(i)}(\boldsymbol{\theta})$$

- This strategy is called Follow the Leader (FTL).

Follow the Reader

- However, there are cases where FTK doesn't work.
- Consider action $\theta \in [-1, 1]$ and cost function $f^{(t)}(\theta) = (1/2)(-1)^t\theta$.
- In this case, the action goes back and forth between -1 and 1 except for the first, as $\theta^{(1)} = 0, \theta^{(2)} = -1, \theta^{(3)} = 1, \dots$
- The cost function is $1/2$ except for the first, as $f^{(1)}(\theta^{(1)}) = 0, f^{(2)}(\theta^{(2)}) = 1/2, f^{(3)}(\theta^{(3)}) = 1/2, \dots$

Follow the Leader

- On the other hand, The optimal strategy is $\theta = 0$ and the total cost of it is 0.
- Therefore, FTL regret in this case doesn't approach 0.
- We have to expand FTL.

Regularized Follow the Reader

- Regularized Follow the Reader (RFTL)

$$\boldsymbol{\theta}^{(t)} = \arg \min_{\boldsymbol{\theta} \in K} \eta \sum_{i=1}^{t-1} f^{(i)}(\boldsymbol{\theta}) + R(\boldsymbol{\theta})$$

- Let $R(\boldsymbol{\theta})$ be convex regularization function. Let $\eta \geq 0$ be parameter that determines the degree of regularization.
- When choosing the first action, the cost function is not presented, so action is determined only by the regularization term.

$$\boldsymbol{\theta}^{(1)} = \arg \min_{\boldsymbol{\theta} \in K} R(\boldsymbol{\theta})$$